

Welcome to Stream Processing with Apache Kafka



UDACITY

The background of the slide features a light blue map-like pattern. It consists of several wavy, solid blue lines that meander across the frame. Interspersed among these lines are small, solid blue dots and faint, dashed blue lines, some of which end in small 'x' marks, suggesting a network or a series of points of interest.

Welcome Message

Stream Processing with Apache Kafka


The background of the slide is a light blue map. It features several solid blue lines that curve across the frame, resembling geographical features like rivers or coastlines. Interspersed among these are dashed blue lines. Small, solid blue dots are placed at various points along the solid lines. Additionally, small blue 'x' marks are scattered across the map, some near the dashed lines and others in open areas.

Understanding Stream Processing

What is Stream Processing?

A **stream** is a potentially unbounded sequence of data


Stream Processing is the act of performing continual calculations on a stream



What is Stream Processing?


Data may be produced at an **even** or **constant** rate

Data may also be produced **unevenly** and in **different shapes and sizes**



What is Stream Processing?


Data streams are made up of **immutable data**. Data cannot be changed once in the stream.



What is Stream Processing?

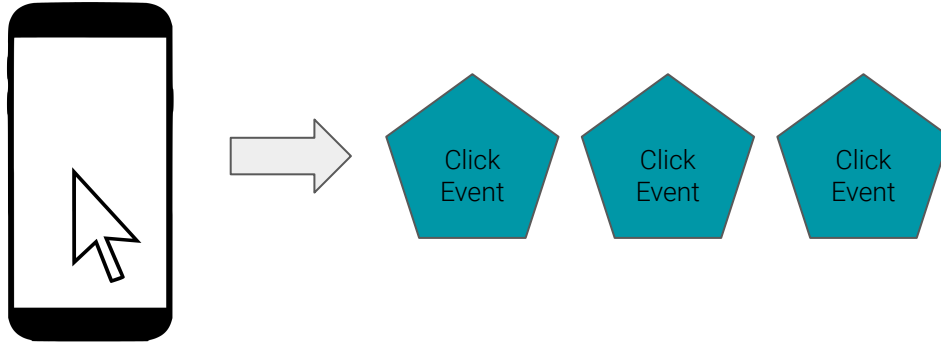
Data records in streams is typically **small**, usually **less than 1MB**

Data **throughput** may range from **one record** per second up to **many thousands** per second




What is an event?

An **event** is an **immutable fact** regarding something that occurred within our system



What is an event?

SQL Databases are built to **store the state** of an application **at that point in time**




user	
id	favorite_color
123	"brown"

What is an event?

SQL Databases are built to **store the state** of an application **at that point in time**

1 UPDATE user
SET favorite_color="red"
WHERE id=123



user	
id	favorite_color
123	"red"

What is an event?

SQL Databases are built to **store the state** of an application **at that point in time**

1

UPDATE user

SET favorite_color="red"


WHERE id=123

2

UPDATE user

SET favorite_color="green"

WHERE id=123



user	
id	favorite_color
123	"green"

What is an event?

SQL Databases are built to **store the state** of an application **at that point in time**

1


```
UPDATE user  
SET favorite_color="red"  
WHERE id=123
```

2

```
UPDATE user  
SET favorite_color="green"  
WHERE id=123
```

3


```
UPDATE user  
SET favorite_color="blue"  
WHERE id=123
```



user	
id	favorite_color
123	"blue"


What is an event?

Traditional **Message Queues** are used to **communicate commands** to **perform an action**



What is an event?

Evented systems **react** to the facts communicated to them. The communication is **indirect** and the events they utilize are usually not **specifically targeted** to any one system.



```
{  
  "action": "click",  
  "element": "search"  
}
```

```
{  
  "action": "click",  
  "element": "view"  
}
```

```
{  
  "action": "click",  
  "element": "checkout"  
}
```

The background of the slide features a light blue map-like pattern. It includes several solid blue lines of varying thicknesses that curve across the frame. Interspersed among these are dashed blue lines, some of which end in small 'x' marks. There are also several small, solid blue dots scattered across the map. A white rectangular box with a thin blue border is centered on the slide, containing the title text.

Examples of Stream Processing


Example: Log Analysis

Logs are hard to process in batch systems due to the speed and size of data

```
2019-10-08 file.go:23 exception | failed opening file
2019-10-08 ormd.go:11 warn      | long transaction
2019-10-08 file.go:70 exception | failed opening file
```

```
2019-10-08 file.go:23 exception | failed opening file
2019-10-08 teld.go:20 info      | dialing 192.168.1.1
2019-10-08 ormd.go:11 warn      | long transaction
2019-10-08 file.go:70 exception | failed opening file
2019-10-08 webp.go:99 info      | opening socket
```

```
2019-10-08 file.go:23 exception | failed opening file
2019-10-08 webp.go:99 info      | opening socket
2019-10-08 file.go:70 exception | failed opening file
```




```
2019-10-08 file.go:23 exception | failed opening file
2019-10-08 teld.go:20 info      | dialing 192.168.1.1
2019-10-08 ormd.go:11 warn      | long transaction
2019-10-08 file.go:70 exception | failed opening file
2019-10-08 webp.go:99 info      | opening socket
```

```
2019-10-08 file.go:23 exception | failed opening file
2019-10-08 webp.go:99 info      | opening socket
2019-10-08 file.go:70 exception | failed opening file
```

```
2019-10-08 file.go:23 exception | failed opening file
2019-10-08 ormd.go:11 warn      | long transaction
2019-10-08 file.go:70 exception | failed opening file
```


Example: Log Analysis

Companies push **log data** as **events** into a **data stream** to perform stream processing.




```
2019-10-08 file.go:23 exception | failed opening file
2019-10-08 teld.go:20 info      | dialing 192.168.1.1
2019-10-08 file.go:23 exception | failed opening file
2019-10-08 webp.go:99 info      | opening socket
2019-10-08 ormd.go:11 warn      | long transaction
2019-10-08 ormd.go:11 warn      | long transaction
2019-10-08 teld.go:20 info      | dialing 10.0.0.1
2019-10-08 teld.go:20 info      | dialing 172.168.0.1
```

Example: Web Analytics

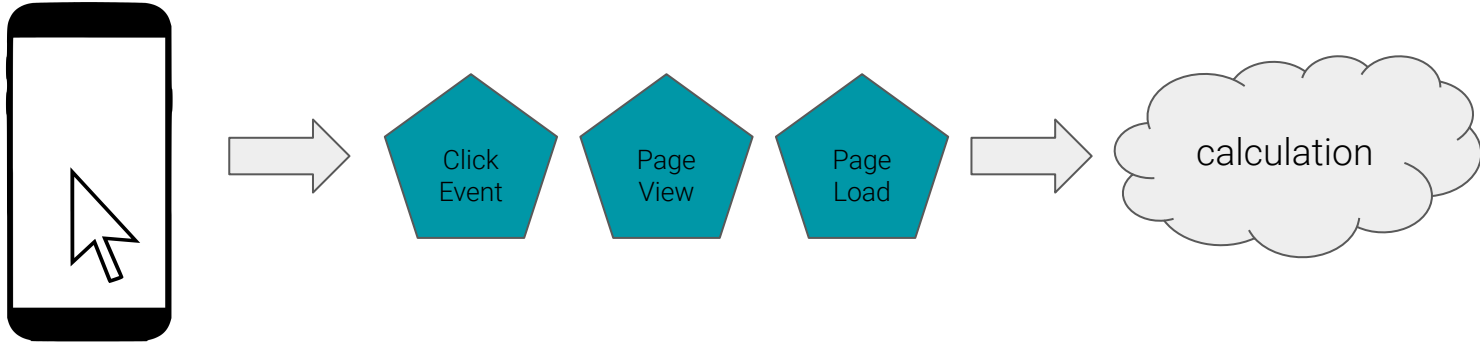
Modern web applications track user actions such as **clicks** and **page views**.

These actions add up fast, overwhelming traditional data stores.



Example: Web Analytics

Stream processing allows companies to process data **as it's generated** and not hours after the fact as is common with batch processing.

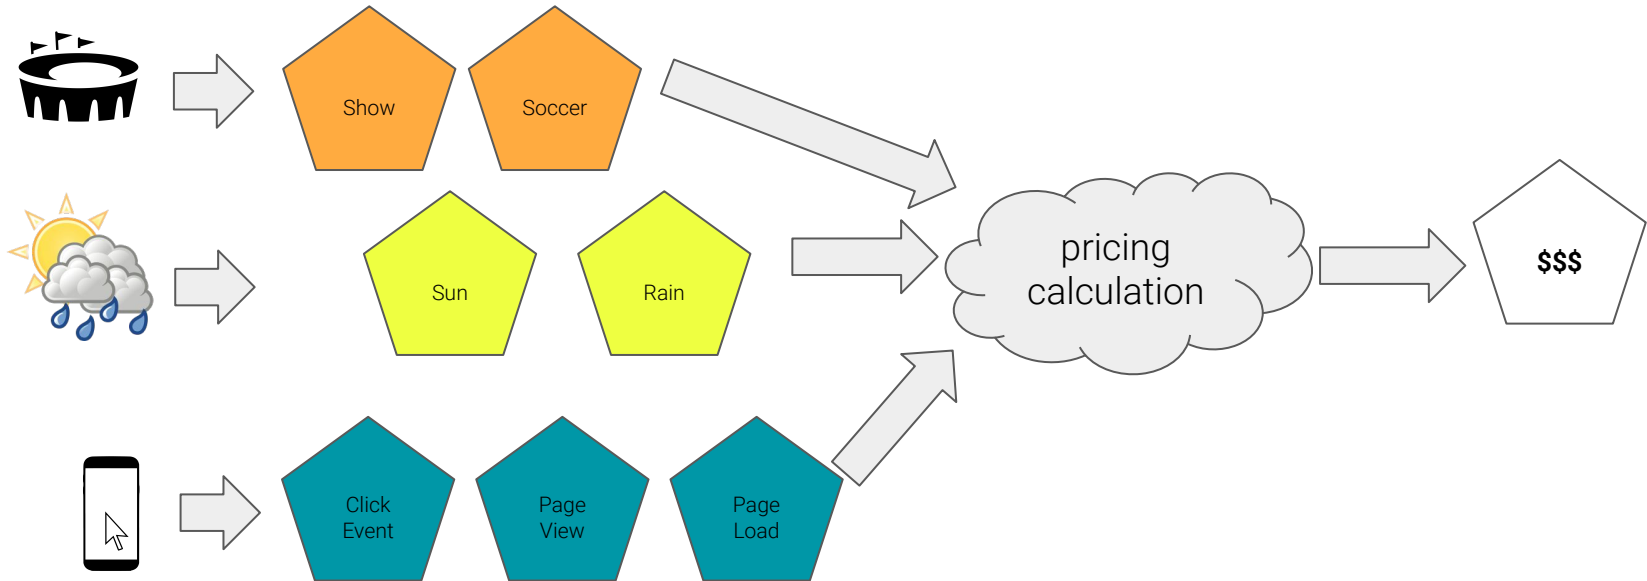


Example: Real-Time Pricing

Real Time Pricing adjusts to **environmental** factors and **instantaneous demand**


Example: Real-Time Pricing

Real Time Pricing adjusts to **environmental** factors and **instantaneous demand**




Example: Financial Analysis

Prices for **stocks fluctuate rapidly**, creating huge amounts of data



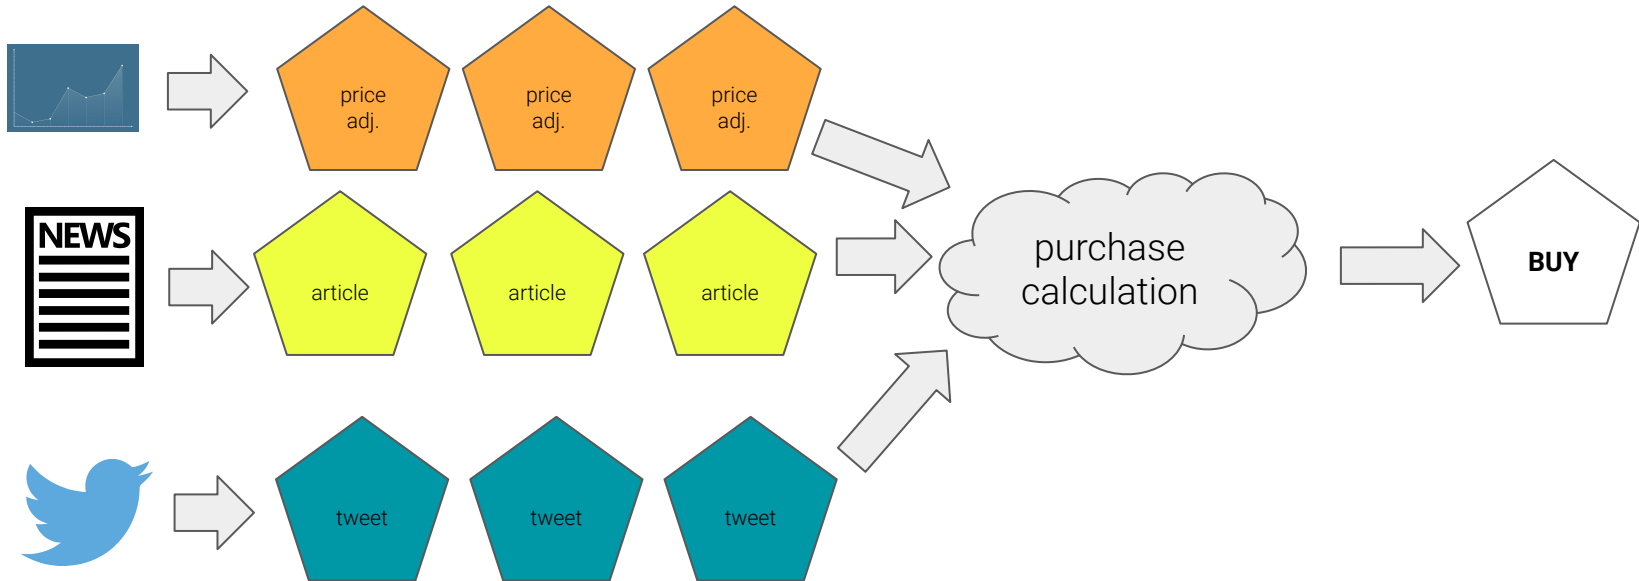
Example: Financial Analysis

Data from many other sources, such as news articles and Twitter influence purchasing




Example: Financial Analysis

Stream Processing coalesces these data streams into one real-time decision to buy or sell




The background of the slide features a light blue map-like pattern. It includes several solid blue lines of varying thicknesses that curve across the frame. Interspersed among these are dashed blue lines, some of which end in small 'x' marks. There are also several solid blue dots scattered across the map. A white rectangular box with a thin blue border is centered on the slide, containing the title text.

Differences Between Stream and Batch Processing

The background of the slide is a light blue map with various wavy lines, some solid and some dashed, and several small blue dots scattered across it. A white rectangular box with a blue border is centered on the slide, containing the title text.

x: Contrasting Stream and Batch Processing

Batch Processing



- Scheduled analysis of related groups of data
- As up to date as the last scheduled run
- May run for long periods of time
- Often involve mutable data stores
- May access all historical data

Stream Processing



- Real-time analysis of data as it is produced
- Runs as soon as event is produced
- As up to date as the last event generated
- Often involves immutable data stores
- Often uses recently produced, windowed data


Batch vs Stream Processing

The differences are **generalizations, not hard rules.**

Most data engineering teams use both **batch** and **stream processing together** to achieve their goals.


Components of a Stream Processing Solution

Stream Processing applications consist of a **streaming data store** and the **streaming calculation(s)**.




Components of a Stream Processing Solution

Streaming data stores **hold all of the immutable event data** in a system. These data stores guarantee that **data is stored in the order it was produced**.




Components of a Stream Processing Solution

Apache **Kafka** and Apache **Cassandra** are examples of streaming data stores.





Components of a Stream Processing Solution

Stream processing applications are **downstream** of the data store and perform calculations on one or more data streams, such as **aggregations**, **joins**, and **filtering**. These calculations **produce new data events**.



Components of a Stream Processing Solution

Popular Stream Processing tools, all of which share common features, include **Confluent KSQL**, **Kafka Streams**, **Faust**, **Spark Structured Streaming**, **Apache Flink**, and **Samza**.



Stream Processing Benefits



Stream Processing Benefits

- **Faster calculations** on windowed data



Stream Processing Benefits

- **Faster calculations** on windowed data
- More **scalable** due to distributed nature of data stores



Stream Processing Benefits

- **Faster calculations** on windowed data
- More **scalable** due to distributed nature of data stores
- Decouples how data is ***used*** from how it is ***produced***



Stream Processing Benefits



- **Faster calculations** on windowed data
- More **scalable** due to distributed nature of data stores
- Decouples how data is ***used*** from how it is ***produced***
- Immutable data provides pipeline **repeatability**

The background of the slide features a light blue map-like pattern. It consists of several wavy, solid blue lines and dashed blue lines. Small blue dots are placed at various points along these lines, and small blue 'x' marks are scattered across the background.

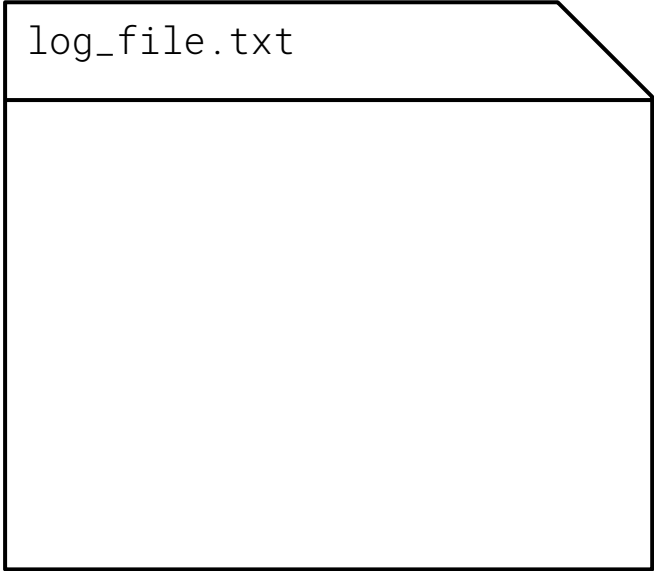
Append-Only Logs

The background of the slide features a light blue map-like pattern. It includes several solid blue lines of varying thicknesses that curve across the frame. Interspersed among these are dashed blue lines, some of which end in small 'x' marks. There are also several small, solid blue dots scattered across the map. A white rectangular box with a thin blue border is centered on the slide, containing the text 'Append-Only Logs'.

Append-Only Logs

Append-Only Logs

Append-only logs are files which **append events to the end as they arrive**



A diagram of a log file. It consists of a rectangular box with a black border. The top-left corner of the box is folded over, creating a triangular flap. The text "log_file.txt" is written in a monospaced font on this flap. The main body of the box is empty, representing the log entries.

log_file.txt

Append-Only Logs

Append-only logs are files which **append events to the end as they arrive**


2019-10-08 **exception** | failed opening file



Append-Only Logs

Append-only logs are files which **append events to the end as they arrive**

2019-10-08 **info** | probing 192.168.1.1




log_file.txt

2019-10-08 **exception** | failed opening file

Append-Only Logs

Append-only logs are files which **append events to the end as they arrive**

2019-10-08 **warn** | closing socket cxn



log_file.txt


2019-10-08 **exception** | failed opening file

2019-10-08 **info** | probing 192.168.1.1

Append-Only Logs

Append-only logs are files which **append events to the end as they arrive**

...



log_file.txt

2019-10-08 **exception** | failed opening file

2019-10-08 **info** | probing 192.168.1.1

2019-10-08 **warn** | closing socket cxn


...

The background of the slide is a light blue map with various wavy lines, some solid and some dashed, and several small blue dots scattered across it. A white rectangular box with a blue border is centered on the slide, containing the title text.

Append-Only Logs in SQL Databases

Append-Only Logs in SQL Databases

SQL Databases **use append-only logs** to communicate and synchronize changes in a process known as **Change Data Capture (CDC)**.



The background of the slide features a light blue map-like pattern. It consists of several wavy, solid blue lines and dashed blue lines. Small blue dots are placed at various points along these lines, and small blue 'x' marks are scattered across the background.

Log-structured storage

Log-structured storage

Characteristics

log_file.txt

2019-10-08 **exception** |
failed opening file

2019-10-08 **info** |
probing 192.168.1.1

2019-10-08 **warn** |
closing socket cxn

. . .

Log-structured storage

Characteristics

- Consist of many **append-only logs** on disk

log_file.txt

2019-10-08 **exception** |
failed opening file

2019-10-08 **info** |
probing 192.168.1.1

2019-10-08 **warn** |
closing socket cxn

. . .

Log-structured storage

Characteristics

- Consist of many **append-only logs** on disk
- Files are periodically **merged**, or **joined together** into one file

log_file.txt

2019-10-08 **exception** |
failed opening file

2019-10-08 **info** |
probing 192.168.1.1

2019-10-08 **warn** |
closing socket cxn

. . .

Log-structured storage

Characteristics

log_file.txt

2019-10-08 **exception** |
failed opening file

2019-10-08 **info** |
probing 192.168.1.1

2019-10-08 **warn** |
closing socket cxn

. . .

- Consist of many **append-only logs** on disk
- Files are periodically **merged**, or **joined together** into one file
- Files are periodically **compacted**, where one or more files is **deleted**, typically **based on age**

Log-structured storage

Characteristics

log_file.txt

2019-10-08 **exception** |
failed opening file

2019-10-08 **info** |
probing 192.168.1.1

2019-10-08 **warn** |
closing socket cxn

. . .

- Consist of many **append-only logs** on disk
- Files are periodically **merged**, or **joined together** into one file
- Files are periodically **compacted**, where one or more files is **deleted**, typically **based on age**
- Use many log files, instead of just one, which **increases speed** and reduces I/O **bottlenecks**


Examples of Log-Structured Storage

Cassandra & HBase



- Cassandra and HBase provide **SQL-like** interfaces
- Use **append-only, log-structured streams**
- Look and act like traditional SQL database to end user
- Clusters may consist of thousands of distributed nodes
- Popular for **batch workloads**

Examples of Log-Structured Storage




- Apache **Kafka** is a **message queue** based on log-structured, append-only storage
- Scales to thousands of distributed nodes
- Popular for **Stream Processing**



Apache Kafka as a Stream Processing Tool


Apache Kafka as a Stream Processing Tool

- Kafka **stores events**, not actions or jobs




Apache Kafka as a Stream Processing Tool

- Kafka **stores events**, not actions or jobs
- Data is **distributed** to multiple nodes **by default**




Apache Kafka as a Stream Processing Tool




- Kafka **stores events**, not actions or jobs
- Data is **distributed** to multiple nodes **by default**
- **Highly scalable** with **strict** data **ordering guarantees**

Apache Kafka as a Stream Processing Tool




- Kafka **stores events**, not actions or jobs
- Data is **distributed** to multiple nodes **by default**
- **Highly scalable** with **strict** data **ordering guarantees**
- **Fault-tolerant** to node loss due to distributed nature

Apache Kafka as a Stream Processing Tool



- Kafka **stores events**, not actions or jobs
- Data is **distributed** to multiple nodes **by default**
- **Highly scalable** with **strict** data **ordering guarantees**
- **Fault-tolerant** to node loss due to distributed nature
- Created at **LinkedIn**, now maintained by **Confluent** as an open source product

Apache Kafka as a Stream Processing Tool



- Kafka **stores events**, not actions or jobs
- Data is **distributed** to multiple nodes **by default**
- **Highly scalable** with **strict** data **ordering guarantees**
- **Fault-tolerant** to node loss due to distributed nature
- Created at **LinkedIn**, now maintained by **Confluent** as an open source product
- Streaming data store for **Flink**, **Spark**, and **Samza**

The background of the slide features a light blue pattern of wavy lines, some solid and some dashed, with small blue dots and 'x' marks scattered throughout, resembling a stylized map or network.


Industry Use-Cases

Uber, Netflix, and Segment

Industry Use-Cases

Uber uses Kafka extensively in their **real-time pricing** pipeline.


Uber



Industry Use-Cases

Segment uses Kafka as their **routing backbone** for analytics events


Segment



Industry Use-Cases

As part of Netflix's **Keystone**, Kafka handles billions of events a day

NETFLIX



The background of the slide features a light blue pattern of wavy lines, some solid and some dashed, with small blue dots and 'x' marks scattered throughout, resembling a stylized map or network.

Kafka in Action

Topics, Consumers, and Producers

Kafka Topics

my_first_topic

```
{"event": "click"}
```

```
{"event": "scroll"}
```


```
{"event": "view"}
```

. . .


- Kafka Topics are roughly analogous to SQL Tables
- Topics **organize related events**
- Topics are **not queryable** like SQL tables
- Topics consist of append-only logs
- All **data** in a topic is in **key-value form**
- Many configuration options and considerations

Kafka Producers

- Kafka **Producers** send data to a topic
- Producers are often built with a client library
- Events may be sent one at a time, or in batches
- Many configuration options and considerations



Kafka Consumers



- **Consumers** retrieve events from Kafka Topic(s)
- May consume one or more topics at a time
- Keep track of what events have already been seen
- Many configuration options and considerations

The background of the slide features a light blue map-like pattern with wavy lines and small blue dots, suggesting a geographical or network theme.

Demo: Using Kafka CLI (Command Line Interface) tools to Produce and Consume

The background of the slide is a light blue map with various wavy lines, some solid and some dashed, and several small blue dots scattered across it. A white rectangular box with a blue border is centered on the slide, containing the title text.

Demo: Using Python to Produce and Consume Events